

Chapter 15. Parallel Algorithms


slides are freely borrowed from the web

Computation Models

- Goal of computation model is to provide a realistic representation of the costs of programming.
- Model provides algorithm designers and programmers a measure of algorithm complexity which helps them decide what is “good” (i.e., efficient)

Goal for Modeling

- We want to develop computational models which accurately represent the cost and performance of programs
- If model is poor, optimum in model may not coincide with optimum observed in practice



Models of Computation

What's a model good for??

- Provides a way to think about computers.
Influences design of:
 - Architectures
 - Languages
 - Algorithms
- Provides a way of estimating how well a program will perform.
Cost in model should be roughly same as cost of executing program

The Random Access Machine Model

RAM model of serial computers:

- Memory is a sequence of words, each capable of containing an integer.
- Each memory access takes one unit of time
- Basic operations (add, multiply, compare) take one unit time.
- Instructions are not modifiable
- Read-only input tape, write-only output tape

What about parallel computers

- RAM model is generally considered a very successful “bridging model” between programmer and hardware.
- “Since RAM is so successful, let's generalize it for parallel computers ...”

PRAM [Parallel Random Access Machine]


(Introduced by Fortune and Wyllie, 1978)

PRAM composed of:

- P processors, each with its own unmodifiable program.
- A single shared memory composed of a sequence of words, each capable of containing an arbitrary integer.
- a read-only input tape.
- a write-only output tape.

PRAM model is a synchronous, MIMD, shared address space parallel computer.

PRAM model of computation



- p processors, each with local memory
- Synchronous operation
- Shared memory reads and writes
- Each processor has unique id in range $[1..p]$
- At each unit of time, a processor is either active or idle (depending on id)

Common Simplifying Assumptions

- Infinite number of processors
- Any memory location is uniformly accessible from any processor
- Infinite amount of shared memory
- SIMD (Single Instruction Multiple Data)
 - All processors execute same program
 - At each time step, all processors execute same instruction on different data ("data-parallel")

Variants of PRAM model

	Exclusive Write	Concurrent Write
Exclusive Read	EREW	ERCW
Concurrent Read	CREW	CRCW

More PRAM taxonomy

- Different protocols can be used for reading and writing shared memory.
 - EREW - exclusive read, exclusive write
A program isn't allowed to have two processors access the same memory location at the same time.
 - CREW - concurrent read, exclusive write
 - ERCW - exclusive read, concurrent write
 - CRCW - concurrent read, concurrent write
Needs protocol for arbitrating write conflicts
- PRAM can emulate a message-passing machine by partitioning memory into private memories.

Why study PRAM algorithms?

- Well-developed body of literature on design and analysis of such algorithms
- Baseline model of concurrency
- Explicit model
 - Specify operations at each step
 - Scheduling of operations on processors
- Robust design paradigm

Work-Span paradigm

- Higher-level abstraction for PRAM algorithms
- Span (Critical Path Length): $T(n)$
 - Most expensive path from beginning to end of a PRAM algorithm
- Work: $W(n)$
 - Total amount of processor time to complete executing the algorithm
- *Work-efficient* if $W(n) = \Theta(T_s(n))$

optimal sequential
algorithm's time complexity

Designing PRAM algorithms

- Balanced trees
- Pointer jumping
- Divide and conquer
- . . .


Balanced trees

- Key idea: Build balanced binary tree on input data, sweep tree up and down
- "Tree" not a data structure, often a control structure (e.g., recursion)

Parallel Sum

Parallel Sum

- Given: Sequence a of $n = 2^k$ elements
- Given: Binary associative operator $+$
- Compute: $S = a_1 + \dots + a_n$



Parallel sum pseudo code

```
integer B[1..n]
forall i in 1 : n do
    B[i] := ai
enddo
for h = 1 to k do
    forall i in 1 : n/2h do
        B[i] := B[2i-1] + B[2i]
    enddo
enddo
S := B[1]
```

Interesting Points

- Global program: no references to processor ID
- Contains both serial and concurrent operations
- Semantics of *forall*
 - In each iteration of the *forall* loop, a processor does the addition or sits idle depending on its ID
- Order of additions different from sequential order: associativity critical

Analysis of parallel sum

- Algorithm is correct
- $\Theta(\lg n)$ steps
- $\Theta(n)$ work
 - In total, $n/2 + n/4 + \dots + 1 = n-1$ additions
- EREW model
- If n not power of 2, pad to next power

Complexity measures of parallel sum

$$T(n) = 1 + k + 1 = \Theta(\lg n)$$

$$W(n) = n + \sum_{h=1}^k \frac{n}{2^h} + 1$$
$$= \Theta(n)$$

- Span is $O(\lg n)$:
Concurrent execution reduces number of steps (shorter critical path length)
- Work: $O(n)$
- Speedup: $O(n/\lg n)$
 - Optimal serial sum algorithm is $O(n)$

Two vertical bars, one red and one yellow, are positioned on the left side of the slide.

Parallel Prefix Sum

What is prefix sum ?

- Input: Sequence x of $n = 2^k$ elements, binary associative operator $+$
- Output: Sequence s of $n = 2^k$ elements, with
$$s_k = x_1 + \dots + x_k$$
- Example:
$$x = [1, 4, 3, 5, 6, 7, 0, 1]$$
$$s = [1, 5, 8, 13, 19, 26, 26, 27]$$

(Inclusive) Prefix-Sum (Scan) Definition

Definition: *The all-prefix-sums operation takes a binary associative operator \oplus , and an array of n elements*

$$[x_0, x_1, \dots, x_{n-1}],$$

and returns the array

$$[x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-1})].$$

Example: If \oplus is addition, then the all-prefix-sums operation on the array $[3 \ 1 \ 7 \ 0 \ 4 \ 1 \ 6 \ 3]$, would return $[3 \ 4 \ 11 \ 11 \ 15 \ 16 \ 22 \ 25]$.

Inclusive Scan Application Example

- Assume we have a 100-inch sandwich to feed 10
- We know how many inches each person wants
 - [3 5 2 7 28 4 3 0 8 1]
- How do we cut the sandwich quickly?
- How much will be left?

- Method 1: cut the sections sequentially: 3 inches first, 5 inches second, 2 inches third, etc.
- Method 2: calculate Prefix scan and cut in parallel
 - [3, 8, 10, 17, 45, 49, 52, 52, 60, 61] (39 inches left)

Typical Applications of Scan

- Scan is a simple and useful parallel building block

- Convert recurrences from **sequential**:

```
for (j=1; j<n; j++)  
    out[j] = out[j-1] + f(j);
```

- into **parallel**:

```
forall(j) { temp[j] = f(j) };  
scan(out, temp);
```

- Useful for many parallel algorithms:

- Radix sort
- Quicksort
- String comparison
- Lexical analysis
- Stream compaction
- Polynomial evaluation
- Solving recurrences
- Tree operations
- Histograms
- Etc.

Other Applications

- Assigning space in farmers market
- Allocating memory to parallel threads
- Allocating memory buffer for communication channels
- ...

An Inclusive Sequential Prefix-Sum

Given a sequence $[x_0, x_1, x_2, \dots]$

Calculate output $[y_0, y_1, y_2, \dots]$

Such that

$$y_0 = x_0$$

$$y_1 = x_0 + x_1$$

$$y_2 = x_0 + x_1 + x_2$$

...

Using a recursive definition

$$y_i = y_{i-1} + x_i$$

A Work Efficient C Implementation

```
y[0] = x[0];  
for (i=1; i < Max_i; i++)  
    y[i] = y[i-1] + x[i];
```

Computationally efficient:

N additions needed for N elements - **$O(N)$**

A Naïve Inclusive Parallel Scan

- Assign one thread to calculate each y element
- Have every thread add up all x elements needed for the y element

$$y_0 = x_0$$

$$y_1 = x_0 + x_1$$

$$y_2 = x_0 + x_1 + x_2$$


Span & Work of this naïve algorithm?

Parallel programming is easy as long as you don't care about performance.


Improving Efficiency

- A common parallel algorithm pattern:
Balanced Trees
 - Build a balanced binary tree on the input data and sweep it to and from the root
 - Tree is not an actual data structure, but a concept to determine what each thread does at each step
- For scan:
 - Traverse down from leaves to root building partial sums at internal nodes in the tree
 - Root holds sum of all leaves
 - Traverse back up the tree building the scan from the partial sums

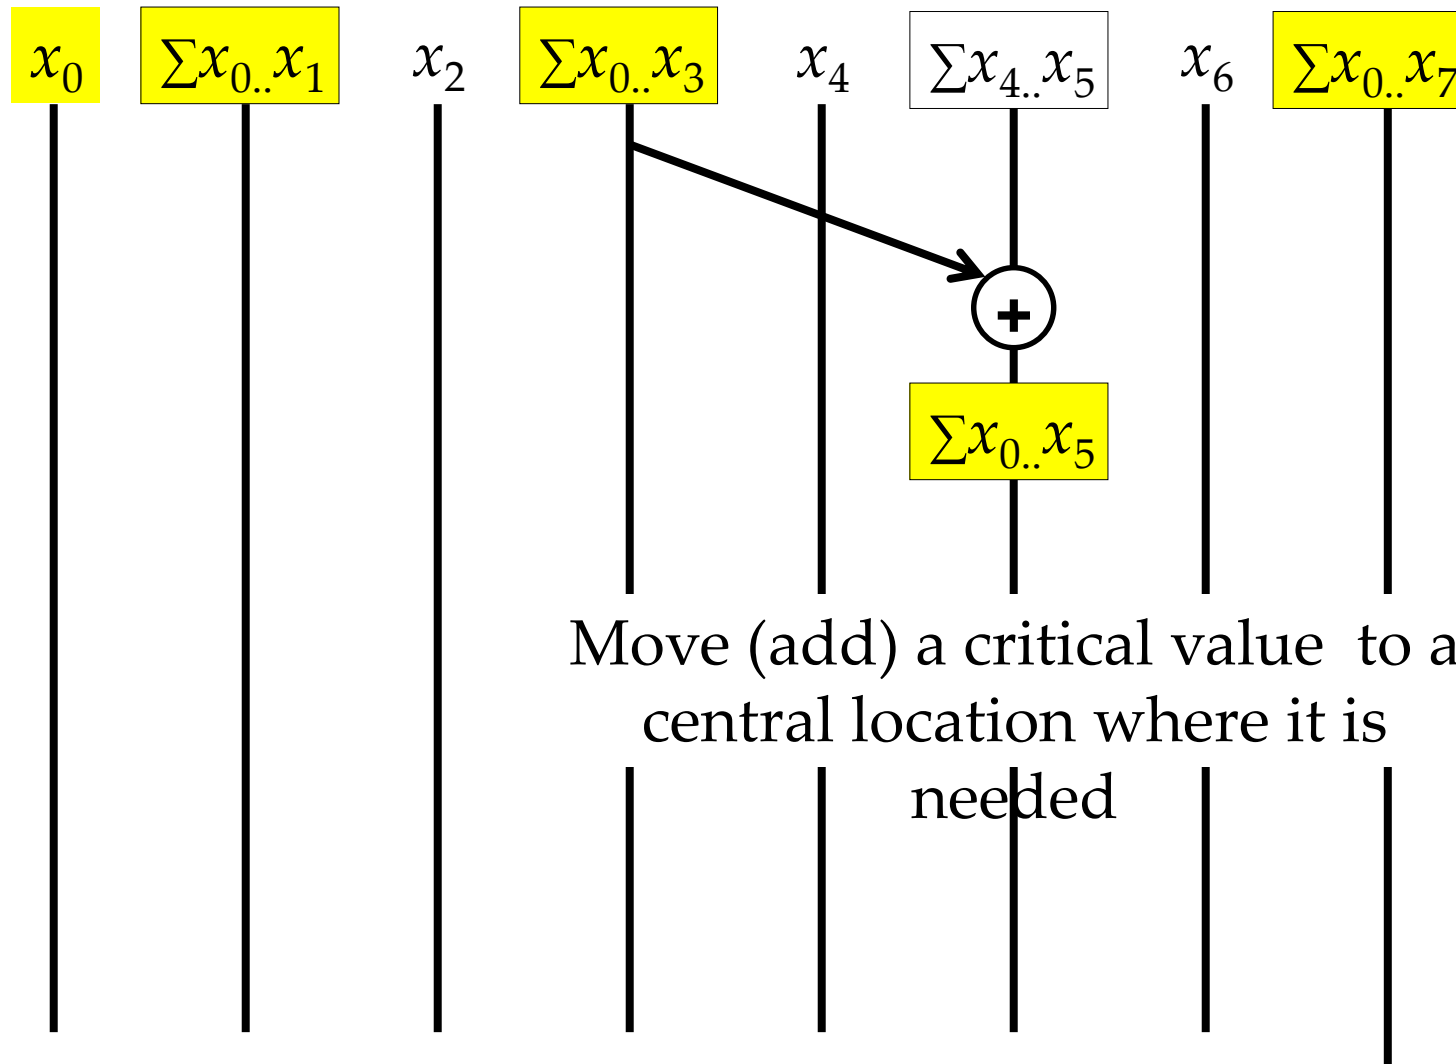
Let's Look at the Reduction Tree Again



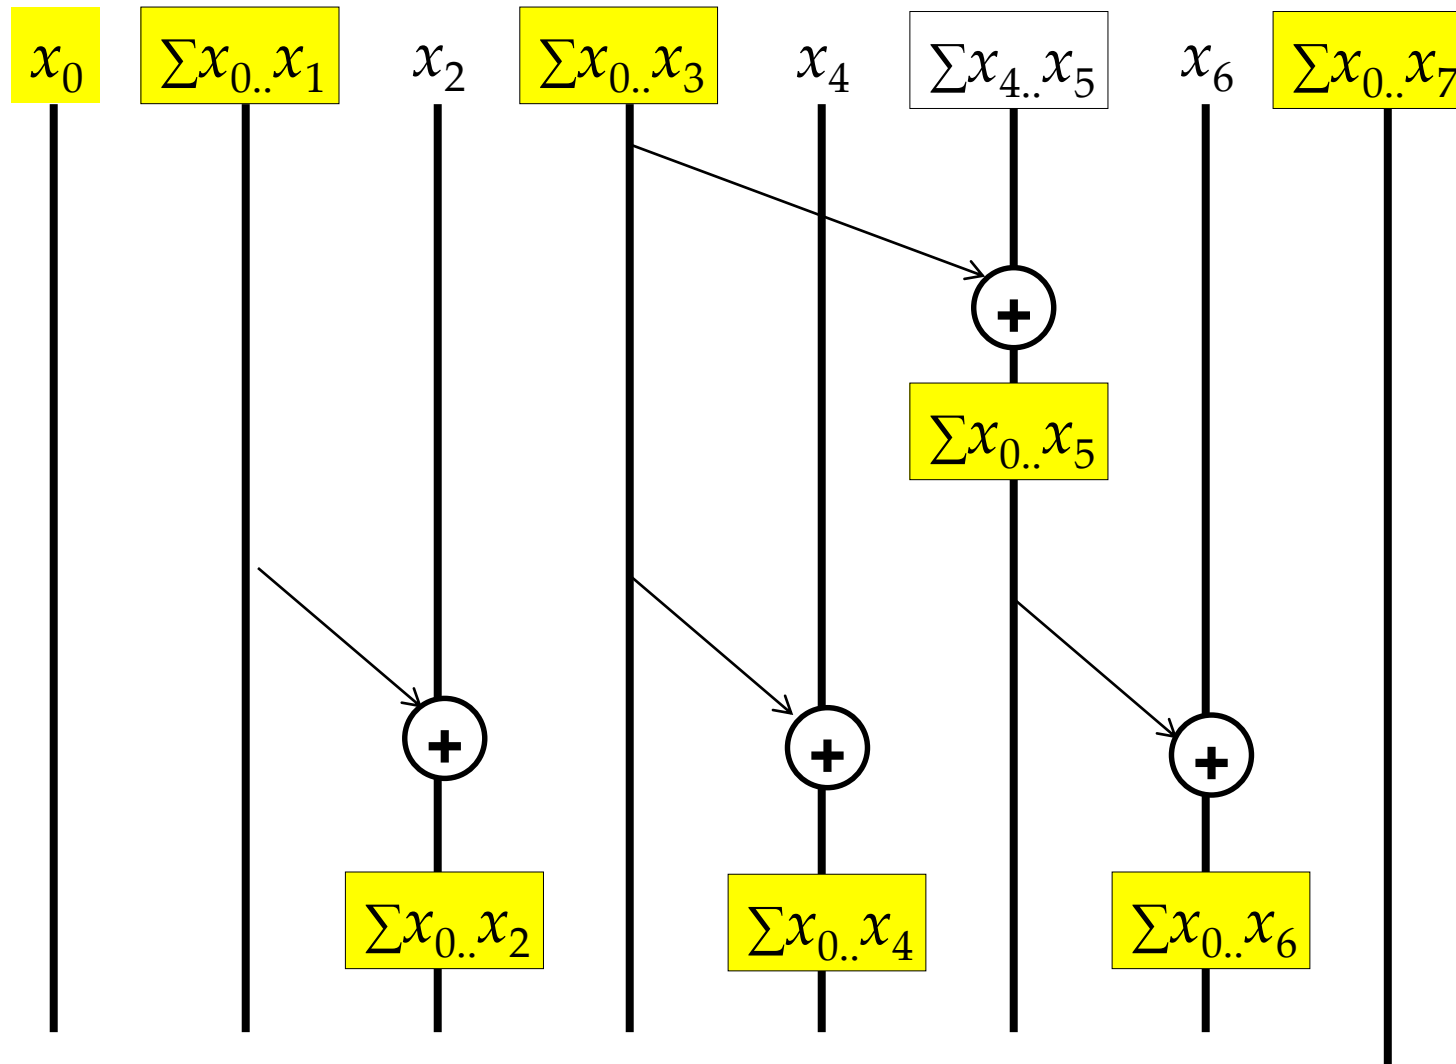
Parallel Scan – Reduction Step




Inclusive Post Scan Step



Inclusive Post Scan Step




Putting it Together



Parallel Prefix Sum Implementations

- There are many multicore CPU implementations of parallel prefix sum
- CUDA implementation:
http://developer.nvidia.com/GPUGems3/gpugems3_ch39.html

Putting it together



(Exclusive) Prefix-Sum (Scan) Definition

Definition: *The all-prefix-sums operation takes a binary associative operator \oplus , and an array of n elements*

$$[x_0, x_1, \dots, x_{n-1}],$$

and returns the array

$$[0, x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-2})].$$

Example: If \oplus is addition, then the all-prefix-sums operation on the array $[3 \ 1 \ 7 \ 0 \ 4 \ 1 \ 6 \ 3]$, would return $[0 \ 3 \ 4 \ 11 \ 11 \ 15 \ 16 \ 22]$.

Why Exclusive Scan


- To find the beginning address of allocated buffers
- Inclusive and Exclusive scans can be easily derived from each other; it is a matter of convenience

[3 1 7 0 4 1 6 3]


Exclusive [0 3 4 11 11 15 16 22]

Inclusive [3 4 11 11 15 16 22 25]

Exclusive Post Scan Step (Add-move Operation)



Exclusive Post Scan Step



Work Analysis

- The parallel Inclusive Scan executes $2 \cdot \log(n)$ parallel iterations. Thus, Span is **$O(\log n)$**
 - $\log(n)$ in reduction and $\log(n)$ in post scan
 - The iterations do $n/2, n/4, \dots, 1, 1, \dots, n/4, n/2$ adds
 - Total adds: $2 \cdot (n-1) \rightarrow \mathbf{O(n)}$ work
- The total number of adds is no more than twice that done in the efficient sequential algorithm
 - The benefit of parallelism can easily overcome the 2X work when there is sufficient hardware

A problem hard to parallelize:

Filter

Filter

[Non-standard terminology]

Given an array `input`, produce an array `output` containing only elements such that `f(elt)` is true

Example: `input [17, 4, 6, 8, 11, 5, 13, 19, 0, 24]`
`f: is elt > 10`
`output [17, 11, 13, 19, 24]`

Looks hard to parallelize

- Finding elements for the output is easy
- But getting them in the right place is hard

Spring 2010

CSE332: Data Abstractions

1

- Slide source: CSE332, University of Washington

Prefix sum to rescue

- $O(\log n)$ span, $O(n)$ work

1. Use a parallel map to compute a **bit-vector** for true elements

input [17, 4, 6, 8, 11, 5, 13, 19, 0, 24]

bits [1, 0, 0, 0, 1, 0, 1, 1, 0, 1]

2. Do parallel-prefix sum on the bit-vector

bitsum [1, 1, 1, 1, 2, 2, 3, 4, 4, 5]

3. Use a parallel map to produce the output

output [17, 11, 13, 19, 24]

```
output = new array of size bitsum[n-1]
if(bitsum[0]==1) output[0] = input[0];
FORALL(i=1; i < input.length; i++)
    if(bitsum[i] > bitsum[i-1])
        output[bitsum[i]-1] = input[i];
```

Parallel quicksort: Algorithm 1

	Best / expected case <i>work</i>
1. Pick a pivot element	$O(1)$
2. Partition all the data into:	$O(n)$
A. The elements less than the pivot	
B. The pivot	
C. The elements greater than the pivot	
3. Recursively sort A and C	$2T(n/2)$

Easy: Do the two recursive calls in parallel

- Work: unchanged of course $O(n \log n)$
- Span: Now $O(n) + 1T(n/2) = O(n)$
- So parallelism (i.e., work/span) is $O(\log n)$

- $O(\log n)$ speedup: Sort 10^9 elements 30 times faster

Parallel quicksort: Algorithm 2

Parallel partition (not in place)

Partition all the data into:

- A. The elements less than the pivot
- B. The pivot
- C. The elements greater than the pivot

- This is just two filters!
 - We know a filter is $O(n)$ work, $O(\log n)$ span
 - Filter elements less than pivot into left side of `aux` array
 - Filter elements great than pivot into right size of `aux` array
 - Put pivot in-between them and recursively sort
 - With a little more cleverness, can do both filters at once but no effect on asymptotic complexity
- With $O(\log n)$ span for partition, the total span for quicksort is $O(\log n) + 1T(n/2) = O(\log^2 n)$
- $O(n/\log n)$ speedup!
- How much is the total work $W(n)$ of this quicksort algorithm?

Parallel List Ranking


List Ranking

- List ranking problem
 - Given a singly linked list L with n objects, for each node, compute the distance to the end of the list
- If d denotes the distance
 - $\text{node.d} = \begin{cases} 0 & \text{if node.next} = \text{nil} \\ \text{node.next.d} + 1 & \text{otherwise} \end{cases}$
- Serial algorithm: $O(n)$
- Parallel algorithm
 - Assign one processor to each node
 - Assume there are as many processors as list objects
 - For each node i , do
 1. $i.d = i.d + i.\text{next}.d$
 2. $i.\text{next} = i.\text{next}.\text{next}$ // pointer jumping

List Ranking via Pointer Jumping

- List_ranking(L)
 1. for each node i , in parallel do
 2. if $i.next = nil$ then $i.d = 0$
 3. else $i.d = 1$
 4. while exists a node i , such that $i.next \neq nil$ do
 5. for each node i , in parallel do
 6. if $i.next \neq nil$ then
 7. $i.d = i.d + i.next.d$ // i updates i itself
 8. $i.next = i.next.next$
- Analysis
 - After a pointer jumping, a list is transformed into two (interleaved) lists
 - After that, four (interleaved) lists
 - Each pointer jumping doubles the number of lists and halves their length
 - After $\lceil \log n \rceil$, all lists contain only one node
 - Total time: $O(\log n)$

List Ranking - Example



List Ranking - Discussion

- Synchronization is important
 - In step 8 ($i.\text{next} = i.\text{next}.\text{next}$), all processors must read right hand side before any processor write left hand side
- The list ranking algorithm is EREW
 - If we assume in step 7 ($i.d = i.d + i.\text{next}.d$) all processors read $i.d$ and then read $i.\text{next}.d$
 - If $j.\text{next} = i$, i and j do not read $i.d$ concurrently
- Work performance
 - performs $O(n \log n)$ work since n processors in $O(\log n)$ time
- Work efficient
 - A PRAM algorithm is work efficient w.r.t another algorithm if two algorithms are within a constant factor
 - Is the link ranking algorithm work-efficient w.r.t the serial algorithm?
 - No, because $O(n \log n)$ versus $O(n)$
- Speedup
 - $S = n / \log n$

Parallel Prefix on a List (4)

- Running time (Span): $O(\log n)$
 - After $\lceil \log n \rceil$, all lists contain only one node
- Work performed: $O(n \log n)$
- Speedup
 - $S = n / \log n$

References

- Joseph JaJa, Introduction to Parallel Algorithms, Addison Wesley
- OpenMP: <http://openmp.org/wp/>
- CUDA Programming Guide:
<http://docs.nvidia.com/cuda/cuda-c-programming-guide/#axzz4eumRgF4w>